

# Lab : 7 : Loan Approval Classification using SVM

**Name : Murali kumar R**

**Roll no : 225229120**

## Step 1 : [Understand Data]

In [35]: `import pandas as pd`

In [36]: `loan=pd.read_csv('train_loan.csv')`  
`loan.head()`

Out[36]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplica
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

In [37]: `loan.shape`

Out[37]: (614, 13)

In [38]: `loan.columns`

Out[38]: Index(['Loan\_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self\_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan\_Amount\_Term', 'Credit\_History', 'Property\_Area', 'Loan\_Status'], dtype='object')

In [39]: `type(loan)`

Out[39]: `pandas.core.frame.DataFrame`

In [40]: `loan.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
Loan_ID          614 non-null object
Gender           601 non-null object
Married          611 non-null object
Dependents       599 non-null object
Education        614 non-null object
Self_Employed    582 non-null object
ApplicantIncome  614 non-null int64
CoapplicantIncome 614 non-null float64
LoanAmount       592 non-null float64
Loan_Amount_Term 600 non-null float64
Credit_History   564 non-null float64
Property_Area     614 non-null object
Loan_Status      614 non-null object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.4+ KB
```

In [41]: `loan.count()`

```
Out[41]: Loan_ID          614
Gender           601
Married          611
Dependents       599
Education        614
Self_Employed    582
ApplicantIncome  614
CoapplicantIncome 614
LoanAmount       592
Loan_Amount_Term 600
Credit_History   564
Property_Area     614
Loan_Status      614
dtype: int64
```

## Step 2 : [Data Cleaning]

In [42]: `# Replace numbers as string by integer`

```
In [43]: def string(x):
          if x == '0':
              return 'bad'
          elif x == '1':
              return 'average'
          elif x == '2':
              return 'good'
          else:
              return 'excellent'
```

```
In [58]: loan['Dependents'] = loan['Dependents'].apply(string)
```

```
In [45]: loan.head()
```

Out[45]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	bad	Graduate	No	5849	
1	LP001003	Male	Yes	average	Graduate	No	4583	
2	LP001005	Male	Yes	bad	Graduate	Yes	3000	
3	LP001006	Male	Yes	bad	Not Graduate	No	2583	
4	LP001008	Male	No	bad	Graduate	No	6000	

```
In [46]: # Missing Categorical Columns
```

```
In [61]: cat_cols=['Gender','Married','Dependents','Education','Credit_History']
cont_cols=['ApplicantIncome','CoapplicantIncome','LoanAmount','Loan_Amount_Term',

loan[cat_cols]=loan[cat_cols].fillna(loan.mode().iloc[0])
loan[cont_cols]=loan[cont_cols].fillna(loan.median().iloc[0])
```

```
In [63]: loan['LoanAmount'].fillna(loan['LoanAmount'].mean(), inplace=True)
```

```
In [55]: data=loan.drop(['Loan_ID'],axis=1)
data.head()
```

Out[55]:


	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	Male	No	bad	Graduate	No	5849	0.0
1	Male	Yes	average	Graduate	No	4583	1508.0
2	Male	Yes	bad	Graduate	Yes	3000	0.0
3	Male	Yes	bad	Not Graduate	No	2583	2358.0
4	Male	No	bad	Graduate	No	6000	0.0

### Step 3 : [OPTIONAL : Exploratory Data Analysis - Who got their loan approved]

```
In [66]: import seaborn as sns
```

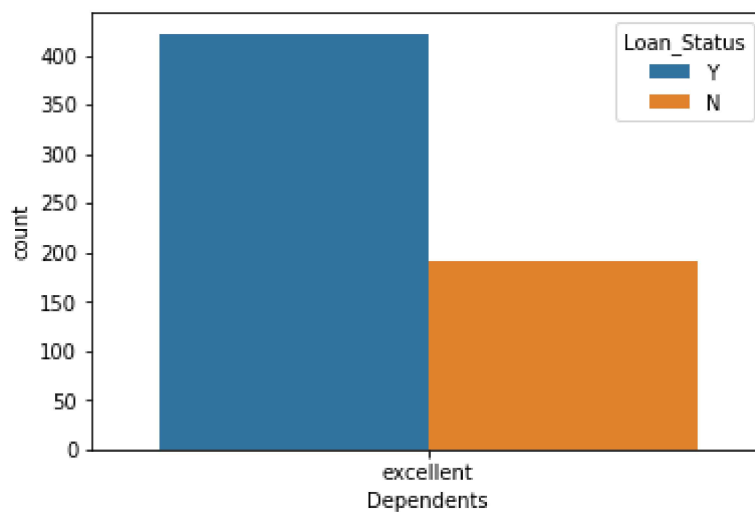
```
In [73]: sns.countplot(x='Married',hue='Loan_Status',data=loan)
print
```

Out[73]: <function print>




```
In [74]: sns.countplot(x='Dependents',hue='Loan_Status',data=loan)
print
```

Out[74]: <function print>




```
In [77]: sns.countplot(x='Education',hue='Loan_Status',data=loan)
print
```

Out[77]: <function print>



```
In [79]: sns.countplot(x='Self_Employed',hue='Loan_Status',data=loan)
print
```

Out[79]: <function print>



## Step 4 : [Extract X and y]

```
In [81]: X = loan.drop(['Loan_Status'],axis=1)
y = loan.Loan_Status
```

```
In [82]: X.head()
```

```
Out[82]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplica
0	LP001002	Male	No	excellent	Graduate	No	5849	
1	LP001003	Male	Yes	excellent	Graduate	No	4583	
2	LP001005	Male	Yes	excellent	Graduate	Yes	3000	
3	LP001006	Male	Yes	excellent	Not Graduate	No	2583	
4	LP001008	Male	No	excellent	Graduate	No	6000	

```
In [83]: y.head()
```

```
Out[83]: 0    Y
1    N
2    Y
3    Y
4    Y
Name: Loan_Status, dtype: object
```

## Step 5 : [One Hot Encoding]

```
In [84]: import warnings
warnings.filterwarnings('ignore')
```

```
In [87]: X=pd.get_dummies(X)
X.head()
```

```
Out[87]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_ID_
0	5849	0.0	3812.5	360.0	1.0	
1	4583	1508.0	128.0	360.0	1.0	
2	3000	0.0	66.0	360.0	1.0	
3	2583	2358.0	120.0	360.0	1.0	
4	6000	0.0	141.0	360.0	1.0	

5 rows × 631 columns

## Step 6 : [Model Building]

```
In [88]: # Spilt X and y for Training and testing
```

```
In [89]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_stat
```

```
In [90]: # Using StandardScaler, fit_transform:
```

```
In [92]: from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
scale
```

```
Out[92]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
In [99]: ss=scale.fit_transform(X_train)
ss1=scale.transform(X_test)

ss
ss1
```

```
Out[99]: array([[ 0.21857952, -0.54770577, -0.29127367, ..., -0.63761852,
 1.29663025, -0.71453718],
 [-0.20156346, -0.54770577, -0.25238992, ..., -0.63761852,
 1.29663025, -0.71453718],
 [ 0.07848281,  0.79027582,  0.04201562, ..., -0.63761852,
 -0.77122989,  1.3995073 ],
 ...,
 [-0.35814214, -0.54770577, -0.35793153, ..., -0.63761852,
 1.29663025, -0.71453718],
 [-0.34578066, -0.54770577, -0.17184501, ...,  1.56833588,
 -0.77122989, -0.71453718],
 [-0.10370178, -0.54770577, -0.26627697, ..., -0.63761852,
 1.29663025, -0.71453718]])
```

```
In [93]: # Linear SVC model:
```

```
In [103]: from sklearn.svm import LinearSVC
model = LinearSVC()
model.fit(ss,y_train)
```

```
Out[103]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
 intercept_scaling=1, loss='squared_hinge', max_iter=1000,
 multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
 verbose=0)
```

```
In [104]: Lsvc_y_pred = model.predict(ss1)
Lsvc_y_pred
```

```
Out[104]: array(['Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y',
                'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y',
                'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'N', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'Y',
                'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N',
                'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'N', 'Y', 'Y',
                'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N',
                'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
                'Y', 'Y', 'Y'], dtype=object)
```

```
In [102]: # Print Accuracy value:
```

```
In [105]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
accuracy_score(y_test, Lsvc_y_pred)
```

```
Out[105]: 0.8324324324324325
```

```
In [106]: # Print Confusion matrix:
```

```
In [109]: confusion_matrix(y_test, Lsvc_y_pred)
```

```
Out[109]: array([[ 22,  29],
                [  2, 132]], dtype=int64)
```

```
In [107]: #print Classification_report:
```

```
In [108]: print(classification_report(y_test, Lsvc_y_pred))
```

	precision	recall	f1-score	support
N	0.92	0.43	0.59	51
Y	0.82	0.99	0.89	134
avg / total	0.85	0.83	0.81	185

## Step 7 : [Performance Comparisons]

```
In [110]: # Compare the performance of LinearSVC against LogisticRegression:
```



```
In [111]: from sklearn.linear_model import LogisticRegression
lgr= LogisticRegression()
lgr.fit(ss,y_train)
lr_y_pred = lgr.predict(ss1)
from sklearn.svm import LinearSVC
l_svc = LinearSVC()
l_svc.fit(ss,y_train)
lsvc_y_pred = l_svc.predict(ss1)
print("LogisticRegression:",accuracy_score(y_test,lr_y_pred))
print("LinearSVC :",accuracy_score(y_test,lsvc_y_pred))
```

LogisticRegression: 0.8324324324324325  
LinearSVC : 0.8324324324324325

```
In [112]: # Compare the performance of LinearSVC against SGDClassifier:
```

```
In [113]: from sklearn.linear_model import SGDClassifier
sgd = SGDClassifier()
sgd.fit(ss,y_train)
sgdc_y_pred = sgd.predict(ss1)
from sklearn.svm import LinearSVC
l_svc = LinearSVC()
l_svc.fit(ss,y_train)
lsvc_y_pred = l_svc.predict(ss1)
print("SGDClassifier:", accuracy_score(y_test,sgdc_y_pred))
print("LinearSVC :",accuracy_score(y_test,lsvc_y_pred))
```

SGDClassifier: 0.8324324324324325  
LinearSVC : 0.8324324324324325

```
In [119]:
```

```
In [124]: from sklearn.svm import SVC

l_svc = LinearSVC()
l_svc.fit(ss,y_train)
lsvc_y_pred = l_svc.predict(ss1)

poly_svc = SVC(kernel='poly', C = 1.0)
poly_svc.fit(ss,y_train)
psvc_y_pred=poly_svc.predict(ss1)

rbf_svc = SVC(kernel='rbf', C = 1.0)
rbf_svc.fit(ss,y_train)
rbfsvc_y_pred=rbf_svc.predict(ss1)

sigmoid_svc = SVC(kernel='sigmoid', C = 1.0)
sigmoid_svc.fit(ss,y_train)
sigsvc_y_pred=sigmoid_svc.predict(ss1)

print("LinearSVC :",accuracy_score(y_test,lsvc_y_pred))
print("poly SVC :",accuracy_score(y_test,psvc_y_pred))
print("rbf SVC :",accuracy_score(y_test,rbfsvc_y_pred))
print("Sigmoid SVC :",accuracy_score(y_test,sigsvc_y_pred))

LinearSVC : 0.8324324324324325
poly SVC : 0.7243243243243244
rbf SVC : 0.7243243243243244
Sigmoid SVC : 0.7621621621621621
```

```
In [125]: # Interpret the results
```

```

In [127]: import numpy as np
from sklearn.metrics import roc_curve
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import auc
MLA = [model,lgr,sgd,poly_svc,rbf_svc,sigmoid_svc]
MLA_columns = []
MLA_compare = pd.DataFrame(columns = MLA_columns)
row_index = 0
for alg in MLA:
    predicted = alg.fit(ss, y_train).predict(ss1)
    predicted=np.where(predicted=='Y',1,0)
    y_testb=np.where(y_test=='Y',1,0)
    fp, tp, th = roc_curve(y_testb, predicted)
    MLA_name = alg.__class__.__name__
    MLA_compare.loc[row_index, 'MLA used'] = MLA_name
    MLA_compare.loc[row_index, 'Train Accuracy'] = round(alg.score(ss,y_train), 4)
    MLA_compare.loc[row_index, 'Test Accuracy'] = round(alg.score(ss1,y_test), 4)
    MLA_compare.loc[row_index, 'Precision'] = precision_score(y_testb, predicted)
    MLA_compare.loc[row_index, 'Recall'] = recall_score(y_testb, predicted)
    MLA_compare.loc[row_index, 'AUC'] = auc(fp, tp)
row_index+=1
MLA_compare

```

Out[127]:

	MLA used	Train Accuracy	Test Accuracy	Precision	Recall	AUC
0	SVC	0.8415	0.7622	0.755682	0.992537	0.5747

In [ ]: