

Lab 10: Patients Physical Activities prediction using Boosting

Name : Murali kumar R

Roll no : 225229120

Step 1: Understand Data

```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
from sklearn.metrics import precision_score, recall_score, accuracy_score, roc_auc_score
from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegressionCV
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score
```


```
In [2]: df = pd.read_csv("Human_Activity_Data.csv")
```

```
In [3]: df.head()
```

Out[3]:

	tBodyAcc- mean()-X	tBodyAcc- mean()-Y	tBodyAcc- mean()-Z	tBodyAcc- std()-X	tBodyAcc- std()-Y	tBodyAcc- std()-Z	tBodyAcc- mad()-X	tBodyAcc- mad()-Y
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111	-0.913526	-0.995112	-0.983185
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300	-0.960322	-0.998807	-0.974914
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187	-0.978944	-0.996520	-0.963668
3	0.279174	-0.026201	-0.123283	-0.996091	-0.983403	-0.990675	-0.997099	-0.982750
4	0.276629	-0.016570	-0.115362	-0.998139	-0.980817	-0.990482	-0.998321	-0.979672

5 rows × 562 columns



```
In [4]: df.shape
```

Out[4]: (10299, 562)

In [5]: `df.columns`

```
Out[5]: Index(['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z',
              'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z',
              'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z',
              'tBodyAcc-max()-X',
              ...,
              'fBodyBodyGyroJerkMag-skewness()', 'fBodyBodyGyroJerkMag-kurtosis()',
              'angle(tBodyAccMean,gravity)', 'angle(tBodyAccJerkMean,gravityMean)',
              'angle(tBodyGyroMean,gravityMean)',
              'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)',
              'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity'],
              dtype='object', length=562)
```

In [6]: `df.dtypes`

```
Out[6]: tBodyAcc-mean()-X          float64
        tBodyAcc-mean()-Y          float64
        tBodyAcc-mean()-Z          float64
        tBodyAcc-std()-X           float64
        tBodyAcc-std()-Y           float64
        ...
        angle(tBodyGyroJerkMean,gravityMean) float64
        angle(X,gravityMean)           float64
        angle(Y,gravityMean)           float64
        angle(Z,gravityMean)           float64
        Activity                      object
        Length: 562, dtype: object
```

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10299 entries, 0 to 10298
Columns: 562 entries, tBodyAcc-mean()-X to Activity
dtypes: float64(561), object(1)
memory usage: 44.2+ MB
```

In [8]: `df["Activity"].value_counts()`

```
Out[8]: LAYING          1944
        STANDING        1906
        SITTING         1777
        WALKING          1722
        WALKING_UPSTAIRS 1544
        WALKING_DOWNSTAIRS 1406
        Name: Activity, dtype: int64
```

Step 2: Build a small dataset

```
In [9]: lay = df.loc[df['Activity'] == "LAYING"][:500]
sit = df.loc[df['Activity'] == "SITTING"][:500]
walk = df.loc[df['Activity'] == "WALKING"][:500]
frames = [lay, sit, walk]
df_new = pd.concat(frames)
```

```
In [10]: df_new.shape
```

```
Out[10]: (1500, 562)
```

Store this dataframe as a new csv file

```
In [11]: df_new.to_csv("Human_Activity_sample.csv")
```

Step 3: Build GradientBoostingClassifier

Import your reduced csv file

```
In [12]: df1= pd.read_csv('Human_Activity_sample.csv')
```


Print Basic Properties of new csv file

```
In [13]: df1.head()
```

```
Out[13]:
```

	Unnamed: 0	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tBodyAcc-mad()-Y	tBodyAcc-mad()-Z
0	51	0.403474	-0.015074	-0.118167	-0.914811	-0.895231	-0.891748	-0.917696	-0.917696	-0.917696
1	52	0.278373	-0.020561	-0.096825	-0.984883	-0.991118	-0.982112	-0.987985	-0.987985	-0.987985
2	53	0.276555	-0.017869	-0.107621	-0.994195	-0.996372	-0.995615	-0.994901	-0.994901	-0.994901
3	54	0.279575	-0.017276	-0.109481	-0.996135	-0.995812	-0.998689	-0.996393	-0.996393	-0.996393
4	55	0.276527	-0.016819	-0.107983	-0.996775	-0.997256	-0.995422	-0.997167	-0.997167	-0.997167

5 rows × 563 columns



```
In [14]: df1.shape
```

```
Out[14]: (1500, 563)
```

```
In [15]: df1.columns
```

```
Out[15]: Index(['Unnamed: 0', 'tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y',  
               'tBodyAcc-mean()-Z', 'tBodyAcc-std()-X', 'tBodyAcc-std()-Y',  
               'tBodyAcc-std()-Z', 'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y',  
               'tBodyAcc-mad()-Z',  
               ...  
               'fBodyBodyGyroJerkMag-skewness()', 'fBodyBodyGyroJerkMag-kurtosis()',  
               'angle(tBodyAccMean,gravity)', 'angle(tBodyAccJerkMean,gravityMean)',  
               'angle(tBodyGyroMean,gravityMean)',  
               'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)',  
               'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity'],  
              dtype='object', length=563)
```

```
In [16]: df1.dtypes
```

```
Out[16]: Unnamed: 0                int64  
         tBodyAcc-mean()-X         float64  
         tBodyAcc-mean()-Y         float64  
         tBodyAcc-mean()-Z         float64  
         tBodyAcc-std()-X         float64  
         ...  
         angle(tBodyGyroJerkMean,gravityMean) float64  
         angle(X,gravityMean)             float64  
         angle(Y,gravityMean)             float64  
         angle(Z,gravityMean)             float64  
         Activity                        object  
         Length: 563, dtype: object
```

```
In [17]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1500 entries, 0 to 1499  
Columns: 563 entries, Unnamed: 0 to Activity  
dtypes: float64(561), int64(1), object(1)  
memory usage: 6.4+ MB
```

```
In [18]: df1["Activity"].value_counts()
```


```
Out[18]: LAYING      500  
         SITTING    500  
         WALKING     500  
         Name: Activity, dtype: int64
```

```
In [19]: X=df1.drop('Activity',axis=1)
X
```

Out[19]:

	Unnamed: 0	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAccmad()-X
0	51	0.403474	-0.015074	-0.118167	-0.914811	-0.895231	-0.891748	-0.917690
1	52	0.278373	-0.020561	-0.096825	-0.984883	-0.991118	-0.982112	-0.987981
2	53	0.276555	-0.017869	-0.107621	-0.994195	-0.996372	-0.995615	-0.994901
3	54	0.279575	-0.017276	-0.109481	-0.996135	-0.995812	-0.998689	-0.996391
4	55	0.276527	-0.016819	-0.107983	-0.996775	-0.997256	-0.995422	-0.997161
...
1495	2637	0.224104	-0.070152	-0.160533	-0.345152	0.143162	-0.430728	-0.381571
1496	2638	0.218667	-0.000151	-0.129555	-0.358372	0.044663	-0.489753	-0.398831
1497	2639	0.282650	0.010504	-0.105785	-0.430536	-0.031534	-0.515552	-0.448291
1498	2640	0.310453	-0.005621	-0.116935	-0.439045	-0.092861	-0.498774	-0.467581
1499	2641	0.350868	-0.029195	-0.120600	-0.367631	0.025890	-0.428028	-0.415911

1500 rows × 562 columns



```
In [20]: y=df1.Activity
y
```

Out[20]:

0	LAYING
1	LAYING
2	LAYING
3	LAYING
4	LAYING
...	...
1495	WALKING
1496	WALKING
1497	WALKING
1498	WALKING
1499	WALKING

Name: Activity, Length: 1500, dtype: object

Split into training and testing set

```
In [21]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=42)
```

Create GradientBoostingClassifier, fit and predict

```
In [22]: model = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1)
model.fit(X_train, y_train)
```

```
Out[22]: GradientBoostingClassifier(learning_rate=1.0, max_depth=1, random_state=42)
```

```
In [23]: y_pred=model.predict(X_test)
```

Print accuracy and classification report

```
In [24]: accuracy_score(y_test, y_pred)
```

```
Out[24]: 1.0
```

```
In [25]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	1.00	1.00	141
WALKING	1.00	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

Step4. [Find Best no. of trees and Best Learning Rate using Grid Search and Cross Validation]

```
In [26]: classifier = GradientBoostingClassifier()
```

```
In [27]: all_scores = cross_val_score(estimator=classifier, X=X_train, y=y_train, cv=5)
```

```
In [28]: all_scores
```

```
Out[28]: array([1., 1., 1., 1., 1.])
```

To find the average of all the accuracies, simple use the mean() method

```
In [29]: all_scores.mean()
```

```
Out[29]: 1.0
```

```
In [30]: parameter = {'n_estimators': [50, 100, 200, 400], 'learning_rate': [0.1, 0.01]}
```

```
In [31]: model1 = GridSearchCV(estimator=classifier, param_grid=parameter, cv=5, n_jobs=
```

```
In [32]: model1.fit(X_train, y_train)
```

```
Out[32]: GridSearchCV(cv=5, estimator=GradientBoostingClassifier(), n_jobs=-1,
               param_grid={'learning_rate': [0.1, 0.01],
                           'n_estimators': [50, 100, 200, 400]})
```

```
In [33]: y_pred2=model1.predict(X_test)
```

```
In [34]: accuracy_score(y_test, y_pred2)
```

```
Out[34]: 1.0
```

```
In [35]: print(classification_report(y_test, y_pred2))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	1.00	1.00	141
WALKING	1.00	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

```
In [36]: print(model1.best_estimator_)
```

```
GradientBoostingClassifier(n_estimators=50)
```

Step5. [Build AdaBoostClassifier]

```
In [37]: base = DecisionTreeClassifier()
```

```
In [38]: model2 = AdaBoostClassifier(base_estimator=base, random_state=0)
```

```
In [39]: param_grid = {'n_estimators': [100, 150, 200], 'learning_rate': [0.01, 0.001]}
```

```
In [40]: model3 = GridSearchCV(model2, param_grid, cv=5, n_jobs=-1)
```

```
In [42]: model3.fit(X_train,y_train)
```

```
Out[42]: GridSearchCV(cv=5,
                    estimator=AdaBoostClassifier(base_estimator=DecisionTreeClassifi
                    er(),
                                                    random_state=0),
                    n_jobs=-1,
                    param_grid={'learning_rate': [0.01, 0.001],
                                'n_estimators': [100, 150, 200]})
```

```
In [43]: y_pred3=model3.predict(X_test)
```

```
In [44]: accuracy_score(y_test,y_pred3)
```

```
Out[44]: 0.9977777777777778
```

```
In [45]: print(classification_report(y_test,y_pred3))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	0.99	1.00	141
WALKING	0.99	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

```
In [46]: print(model3.best_estimator_)
```

```
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.0
1,
                    n_estimators=100, random_state=0)
```

Step6. [Build LogisticRegressionCV classifier]

```
In [47]: model4 = LogisticRegressionCV(cv=4,Cs=5,penalty='l2')
```

```
In [49]: model4.fit(X_train,y_train)
```

```
Out[49]: LogisticRegressionCV(Cs=5, cv=4)
```

```
In [50]: y_pred4=model4.predict(X_test)
```

```
In [51]: accuracy_score(y_test,y_pred4)
```

```
Out[51]: 0.9977777777777778
```



```
In [52]: print(classification_report(y_test,y_pred4))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	0.99	1.00	141
WALKING	0.99	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

Step 7 [Build VotingClassifier]

```
In [53]: model5=VotingClassifier(estimators=[('lr',model4),('gbc',model1)], voting='hard')
```

```
In [54]: model5.fit(X_train,y_train)
```

```
Out[54]: VotingClassifier(estimators=[('lr', LogisticRegressionCV(Cs=5, cv=4)),
                                      ('gbc',
                                       GridSearchCV(cv=5,
                                                    estimator=GradientBoostingClassifi
er(),
                                                    n_jobs=-1,
                                                    param_grid={'learning_rate': [0.1,
                                                                                   0.0
                                                                                   1],
                                                                                   'n_estimators': [50, 1
                                                                                   00,
                                                                                   200,
                                                                                   40
                                                                                   0]})))]
```

```
In [55]: y_pred5=model5.predict(X_test)
```

```
In [56]: print(classification_report(y_test,y_pred5))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	1.00	1.00	141
WALKING	1.00	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

Step8. [Interpret your results]

In [57]: `print(model1.best_estimator_)`

GradientBoostingClassifier(n_estimators=50)

In [58]: `print(model3.best_estimator_)`

AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.01,
n_estimators=100, random_state=0)

GradientBoostingClassifier

GradientBoostingClassifier(n_estimators=50)

In [59]: `classifierF = GradientBoostingClassifier(n_estimators=50)
all_scoresF = cross_val_score(estimator=classifier, X=X_train, y=y_train, cv=5)
parameter = {'n_estimators': [50, 100, 200, 400], 'learning_rate': [0.1, 0.01]}`

In [60]: `modelGC = GridSearchCV(estimator=classifier, param_grid=parameter, cv=5, n_jobs=-1)`

In [61]: `modelGC.fit(X_train, y_train)`

Out[61]: `GridSearchCV(cv=5, estimator=GradientBoostingClassifier(), n_jobs=-1,
param_grid={'learning_rate': [0.1, 0.01],
'n_estimators': [50, 100, 200, 400]})`

In [62]: `y_predGC=model3.predict(X_test)`

In [63]: `accuracy_score(y_test, y_predGC)`

Out[63]: 0.9977777777777778

In [64]: `print(classification_report(y_test, y_predGC))`

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	0.99	1.00	141
WALKING	0.99	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

AdaBoostClassifier

AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),learning_random_state=0)

```
In [66]: modelABC = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning
```

```
In [67]: param_grid = {'n_estimators': [100, 150, 200], 'learning_rate': [0.01, 0.001]}
```

```
In [68]: modelGSCV = GridSearchCV(modelABC,param_grid,cv=5,n_jobs=-1)
modelGSCV.fit(X_train,y_train)
```

```
Out[68]: GridSearchCV(cv=5,
                      estimator=AdaBoostClassifier(base_estimator=DecisionTreeClassifi
er(),
                                                    learning_rate=0.01, n_estimators=10
0,
                                                    random_state=0),
                      n_jobs=-1,
                      param_grid={'learning_rate': [0.01, 0.001],
                                   'n_estimators': [100, 150, 200]})
```

```
In [69]: y_predGSCV=model3.predict(X_test)
```

```
In [70]: accuracy_score(y_test,y_predGSCV)
```

```
Out[70]: 0.9977777777777778
```

```
In [71]: print(classification_report(y_test,y_predGSCV))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	148
SITTING	1.00	0.99	1.00	141
WALKING	0.99	1.00	1.00	161
accuracy			1.00	450
macro avg	1.00	1.00	1.00	450
weighted avg	1.00	1.00	1.00	450

```
In [ ]:
```